

[Skip to content](#)

[Switch to mobile version](#)

[This repository](#)

[Search](#)

[Pull requests](#)

[Issues](#)

[Marketplace](#)

[Explore](#)

[@ChillBoss](#) [Sign out](#)

0

1 1 [rajitnikhare/MNIST](#)

[Code](#) [Issues](#) 0 [Pull requests](#) 0 [Projects](#) 0 [Insights](#)

[MNIST/MNIST_Softmax.py](#)

753e6c9 on 5 Jun 2018

[@rajitnikhare](#) [rajitnikhare](#) [Add files via upload](#)

254 lines (125 sloc) 5.41 KB

coding: utf-8

In[150]:

#Importing tensorflow

import tensorflow as tf

In[151]:

#Importing MNIST Dataset

from tensorflow.examples.tutorials.mnist import input_data

mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

one hot vectors is 0 in most dimension except for a single which denotes a value. Give a picture example

In[152]:

#Creating placeholders

```
x = tf.placeholder(tf.float32, [None, 784])
```

#placeholder is not a variable. It is a value created for the tensorflow to run computation.

#MNIST data is in 28*28 pixel shape which needs to be flattened into 784 dimension

vector($28*28 = 784$). Thus, a 2-D tensor

#is created with a shape of [None, 784](None means dimension can be of varied length).

explain more about placeholders with pictures and diagram

In[153]:

#Creating weights and biases

```
w = tf.Variable(tf.zeros([784,10]))
```

```
b = tf.Variable(tf.zeros([10]))
```

#Variables here is a value that lives in TensorFlow's computation graph. This can be changed and used by the computation.

#initially w, b are zeros of (784,10) and (10) in dimension because it needs to learn the values so it does

#not matter if the initial values are zeros.

#our model is $y = \text{softmax}(x*w + b)$. Thus, w is of the dimension (784*10). Matrix multiplication of x and w is going

#to result in a vector of shape (10). This can be added to b.

#need a diagram here

In[154]:

our model

```
y = tf.nn.softmax(tf.matmul(x,w) + b)
```

#why this model. Explain more about softmax.

In[155]:

#training our model

```
y_ = tf.placeholder(tf.float32, [None,10])
```

#cross entropy function

```
cross_entropy = tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(y), reduction_indices=[1]))
```

#cross entropy function describes how inefficient our predictions are. It is the loss function.

#what's cross entropy? Why use it here? equation of cross entropy function?

In[]:

```
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

#gradient optimizer tries to move to that direction in the computation graph where the cost or loss is reduced.

#learning rate/ step length is 0.5. It will descend through cross_entropy.

#what's Gradient descent? Explain with diagram? fast.ai approach is good here

```
#init = tf.initialize_all_variables()
```

```
init = tf.global_variables_initializer()
```

Variables are not initialized on their own. The above command is needed to initialize them.

In[157]:

#interactive session:

```
sess = tf.Session()
```

sess.run(init) # reset values to incorrect defaults.

#Tensorflow uses a C++ backend. The connection to the backend is a session. Normally, a computation graph is

```
# created and then the session is launched.
```

```
# In[158]:
```

```
for i in range(100):  
    batch_xs, batch_ys = mnist.train.next_batch(100)  
    sess.run(train_step, feed_dict={x : batch_xs, y_ : batch_ys})
```

```
# In[159]:
```

```
#Evaluating the model:  
prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
```

```
# In[160]:
```

```
accuracy = tf.reduce_mean(tf.cast(prediction, tf.float32))
```

```
# In[161]:
```

```
print(sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels}))
```

```
# In[162]:
```

```
# Multilayer Convolutional Layer:
```

```
#functions to define weights and biases to be used in the computation ahead.
```

```
def weight_variable(shape):  
    initial = tf.truncated_normal(shape, stddev=0.1)  
    return tf.Variable(initial)
```

```
def bias_variable(shape):  
    initial = tf.constant(0.1, shape=shape)  
    return tf.Variable(initial)
```

```
# In[175]:
```

```
def conv2d(x, W):  
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')
```

```
def max_pool_2x2(x):  
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],  
                           strides=[1, 2, 2, 1], padding='SAME')
```

```
# In[177]:
```

```
get_ipython().run_line_magic('pinfo', 'tf.nn.conv2d')
```

```
# In[164]:
```

```
#First Convolutional Layer  
W_conv1 = weight_variable([5, 5, 1, 32])  
b_conv1 = bias_variable([32])
```

```
# In[165]:
```

```
x_image = tf.reshape(x, [-1,28,28,1])
```

```
# In[166]:
```

```
h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)  
h_pool1 = max_pool_2x2(h_conv1)
```

```
# In[167]:
```

```
#Second Convolutional Layer
```

```
W_conv2 = weight_variable([5, 5, 32, 64])
```

```
b_conv2 = bias_variable([64])
```

```
h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
```

```
h_pool2 = max_pool_2x2(h_conv2)
```

```
# In[168]:
```

```
W_fc1 = weight_variable([7 * 7 * 64, 1024])
```

```
b_fc1 = bias_variable([1024])
```

```
h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
```

```
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
```

```
# In[169]:
```

```
#Dropout Layer
```

```
keep_prob = tf.placeholder(tf.float32)
```

```
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
```

```
# In[170]:
```

```
#Readout Layer
```

```
W_fc2 = weight_variable([1024, 10])
```

```
b_fc2 = bias_variable([10])
```

```
y_conv = tf.matmul(h_fc1_drop, W_fc2) + b_fc2
```

```
# In[174]:
```

```
with tf.Session() as sess:
```

```
cross_entropy = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(labels=y_, logits=y_conv))
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
correct_prediction = tf.equal(tf.argmax(y_conv,1), tf.argmax(y_,1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
sess.run(tf.global_variables_initializer())

for i in range(20000):
    batch = mnist.train.next_batch(50)
    if i%100 == 0:
        train_accuracy = accuracy.eval(feed_dict={
            x:batch[0], y_: batch[1], keep_prob: 1.0})
        print("step %d, training accuracy %g"%(i, train_accuracy))
        train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5})

print("test accuracy %g"%accuracy.eval(feed_dict={
    x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))
```

© 2019 GitHub, Inc.

[Terms](#)

[Privacy](#)

[Security](#)

[Status](#)

[Help](#)

[Contact GitHub](#)

[Pricing](#)

[API](#)

[Training](#)

[Blog](#)

[About](#)

Press h to open a hovercard with more details.